

Aviation Data Analysis

Overview

My company is planning to expand into the aircraft industry, and is interested in purchasing and operating aircrafts for both commercial and private enterprises. Since this is a novel business venture for the organization, an analysis of aircraft data has been conducted in order to identify potential risks and acquire useful insights that the company can utilize as it diversifies its portfolio.

Business Understanding

The objective of this analysis is to determine if delving into the aircraft industry is worthwhile for the company, and to identify which types of aircrafts the company should purchase to reduce risks. To achieve this, the following questions were investigated:

1. What is the trend of aircraft accidents and fatalities over the years?
2. At which phases of flight are accidents more likely to occur? Which phases result in more fatalities and which phases result in more damage to the aircraft?
3. Is there a relation between number of engines in an aircraft and fatalities incase of an accident?

Data Understanding, Preparation and Analysis

For this analysis, I used a dataset from the National Transportation Safety Board that includes accident data from 1962 to 2023. The data contains records of civil aviation accidents and selected incidents in the USA and international waters.

Loading the Data

In [1]:

```
# Importing required modules and libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

In [2]:

```
# Loading the dataset into a DataFrame
df = pd.read_csv('Data/Aviation_Data.csv')
df.head()

c:\Users\USER\anaconda3\envs\learn-env\lib\site-packages\IPython\core\interactiveshell.py
:3145: DtypeWarning: Columns (6,7,28) have mixed types.Specify dtype option on import or
set low_memory=False.
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

Out[2]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Coc
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	NaN	NaN	Na
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	NaN	NaN	Na

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Code
2	20061025X01558			1974-08-30					
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	NaN	NaN	Na
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	NaN	NaN	Na

5 rows x 31 columns



In [3]:

```
# Displaying the rows and columns of the data
df.shape
```

Out[3]:

```
(90348, 31)
```

In [4]:

```
# Displaying the column names
df.columns
```

Out[4]:

```
Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
      'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
      'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.of.Engines', 'Engine.Type', 'FAR.Description',
      'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')
```

In [5]:

```
# Formatting column names to strip whitespaces, convert to lowercase and replace '.' with '_'
df.columns = df.columns.str.strip().str.lower().str.replace('.', '_')
```

In [6]:

```
# Displaying general information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90348 entries, 0 to 90347
Data columns (total 31 columns):
#   Column                               Non-Null Count  Dtype
---  -
0   event_id                             88889 non-null  object
1   investigation_type                   90348 non-null  object
2   accident_number                      88889 non-null  object
3   event_date                           88889 non-null  object
4   location                             88837 non-null  object
5   country                             88663 non-null  object
6   latitude                             34382 non-null  object
7   longitude                            34373 non-null  object
8   airport_code                         50249 non-null  object
9   airport_name                         52790 non-null  object
10  injury_severity                      87889 non-null  object
11  aircraft_damage                     85695 non-null  object
12  aircraft_category                   32287 non-null  object
13  registration_number                 87572 non-null  object
14  make                                88826 non-null  object
15  model                               88797 non-null  object
16  amateur_built                       88787 non-null  object
17  number_of_engines                   82805 non-null  float64
```

```

18 engine_type            81812 non-null object
19 far_description        32023 non-null object
20 schedule               12582 non-null object
21 purpose_of_flight      82697 non-null object
22 air_carrier            16648 non-null object
23 total_fatal_injuries    77488 non-null float64
24 total_serious_injuries  76379 non-null float64
25 total_minor_injuries    76956 non-null float64
26 total_uninjured        82977 non-null float64
27 weather_condition      84397 non-null object
28 broad_phase_of_flight  61724 non-null object
29 report_status          82508 non-null object
30 publication_date        73659 non-null object

```

```
dtypes: float64(5), object(26)
```

```
memory usage: 21.4+ MB
```

The dataset contains 90,348 records with 31 columns. There are a number of columns such as `event_date`, `aircraft_damage` and `number_of_engines` have missing values.

For ease of analysis, only the columns of interest will be retained.

```
In [7]:
```

```

# Creating a new DataFrame with only the columns of interest
aviation_data = df.loc[:, ['investigation_type', 'event_date', 'aircraft_damage', 'number_of_engines', 'total_fatal_injuries', 'broad_phase_of_flight']]
aviation_data.head()

```

```
Out[7]:
```

	investigation_type	event_date	aircraft_damage	number_of_engines	total_fatal_injuries	broad_phase_of_flight
0	Accident	1948-10-24	Destroyed	1.0	2.0	Cruise
1	Accident	1962-07-19	Destroyed	1.0	4.0	Unknown
2	Accident	1974-08-30	Destroyed	1.0	3.0	Cruise
3	Accident	1977-06-19	Destroyed	1.0	2.0	Cruise
4	Accident	1979-08-02	Destroyed	NaN	1.0	Approach

Missing Values

Most columns in the `aviation_data` DataFrame have missing values. The percentage of missing values in each column is indicated below.

```
In [8]:
```

```
aviation_data.isna().sum().map(lambda x: f"{round((x/aviation_data.shape[0])*100, 2)}%")
```

```
Out[8]:
```

```

investigation_type      0.0%
event_date              1.61%
aircraft_damage          5.15%
number_of_engines        8.35%
total_fatal_injuries     14.23%
broad_phase_of_flight    31.68%
dtype: object

```

Dropping Rows with Missing Values

Records with missing values in the `event_date` column are displayed below.

```
In [9]:
```

```
aviation_data.loc[aviation_data['event_date'].isna()]
```

Out[9]:

	investigation_type	event_date	aircraft_damage	number_of_engines	total_fatal_injuries	broad_phase_of_flight
64030	25-09-2020	NaN	NaN	NaN	NaN	NaN
64050	25-09-2020	NaN	NaN	NaN	NaN	NaN
64052	25-09-2020	NaN	NaN	NaN	NaN	NaN
64388	25-09-2020	NaN	NaN	NaN	NaN	NaN
64541	25-09-2020	NaN	NaN	NaN	NaN	NaN
...
90004	15-12-2022	NaN	NaN	NaN	NaN	NaN
90010	15-12-2022	NaN	NaN	NaN	NaN	NaN
90031	15-12-2022	NaN	NaN	NaN	NaN	NaN
90090	20-12-2022	NaN	NaN	NaN	NaN	NaN
90097	20-12-2022	NaN	NaN	NaN	NaN	NaN

1459 rows × 6 columns

It seems that most records missing the event date are also missing values in the other columns. In that case, dropping these records will not result in any loss of data.

In [10]:

```
# Dropping rows with missing event date
aviation_data.dropna(subset=['event_date'], inplace=True)
```

Replacing Missing Values with Appropriate Category Name

Below are the categories present in the `aircraft_damage` column.

In [11]:

```
# Displaying unique values in aircraft_damage column
aviation_data['aircraft_damage'].unique()
```

Out[11]:

```
array(['Destroyed', 'Substantial', 'Minor', nan, 'Unknown'], dtype=object)
```

Aside from the missing values in this column, there is another category labelled 'Unknown'. The missing values in `aircraft_damage` column can therefore be replaced with the value 'Unknown'.

In [12]:

```
# Replacing missing values in the column with 'Unknown'
aviation_data['aircraft_damage'].fillna(value='Unknown', inplace=True)
```

The same observation can be made for the `broad_phase_of_flight` column, and therefore the same treatment of missing values is applied.

In [13]:

```
# Displaying unique values in broad_phase_of_flight column
aviation_data['broad_phase_of_flight'].unique()
```

Out[13]:

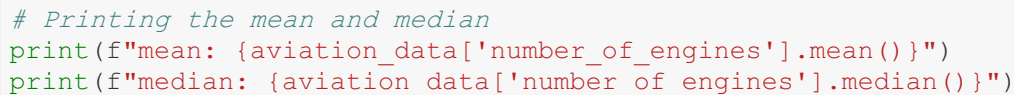
```
array(['Cruise', 'Unknown', 'Approach', 'Climb', 'Takeoff', 'Landing',
      'Taxi', 'Descent', 'Maneuvering', 'Standing', 'Go-around', 'Other',
      nan], dtype=object)
```

```
# Replacing missing values in the column with 'Unknown'
aviation_data['broad phase of flight'].fillna(value='Unknown', inplace=True)
```

Next, the distribution of data in the `number_of_engines` column is observed with the use of a histogram.

```
# Plotting a histogram to view the distribution of number of engines
aviation_data['number_of_engines'].plot(kind='hist', bins=8)


plt.title('Distribution of Number of Engines on an Aircraft')
plt.xlabel('Number of Engines')
plt.ylabel('Aircrafts With This Number of Engines');
```




```
# Imputing missing values with the median
aviation_data['number_of_engines'].fillna(value=aviation_data['number_of_engines'].median(), inplace=True)
```

```
# Plotting a boxplot to view the distribution of total fatal injuries
aviation_data['total_fatal_injuries'].plot(kind='box')

plt.title('Distribution of Total Fatal Injuries');
```





In [19]:

```
# Printing the mean and median
print(f"mean: {aviation_data['total_fatal_injuries'].mean()}")
print(f"median: {aviation_data['total_fatal_injuries'].median()}")
```

```
mean: 0.6478551517654346
median: 0.0
```

Similarly, the distribution of `total_fatal_injuries` is positively skewed. In this case the missing values are also imputed with the median.

In [20]:

```
# Imputing missing values with the median
aviation_data['total_fatal_injuries'].fillna(value=aviation_data['total_fatal_injuries'].
median(), inplace=True)
```

The missing values in our dataset have now been handled.

In [21]:

```
# Displaying percentage of missing values for each column
aviation_data.isna().sum().map(lambda x: f"{round((x/aviation_data.shape[0])*100, 2)}%")
```

Out[21]:

```
investigation_type      0.0%
event_date              0.0%
aircraft_damage         0.0%
number_of_engines       0.0%
total_fatal_injuries    0.0%
broad_phase_of_flight   0.0%
dtype: object
```

As a final step of our data preparation, the datatype of the `event_date` column is changed from object to datetime, which is more appropriate for our analysis.

In [22]:

```
aviation_data['event_date'] = pd.to_datetime(aviation_data['event_date'])
```

In [23]:

```
aviation_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 88889 entries, 0 to 90347
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   investigation_type     88889 non-null  object
1   event_date             88889 non-null  datetime64[ns]
2   aircraft_damage        88889 non-null  object
3   number_of_engines      88889 non-null  float64
4   total_fatal_injuries   88889 non-null  float64
5   broad_phase_of_flight  88889 non-null  object
```

```
dtypes: datetime64[ns](1), float64(2), object(3)
memory usage: 4.7+ MB
```

Now that the data is clean, our data analysis can begin.

Trend of aircraft accidents and fatalities over the years?

To answer our first question, a line graph is used to observe the trend of aircraft accidents and fatal injuries resulting from these accidents over the years.

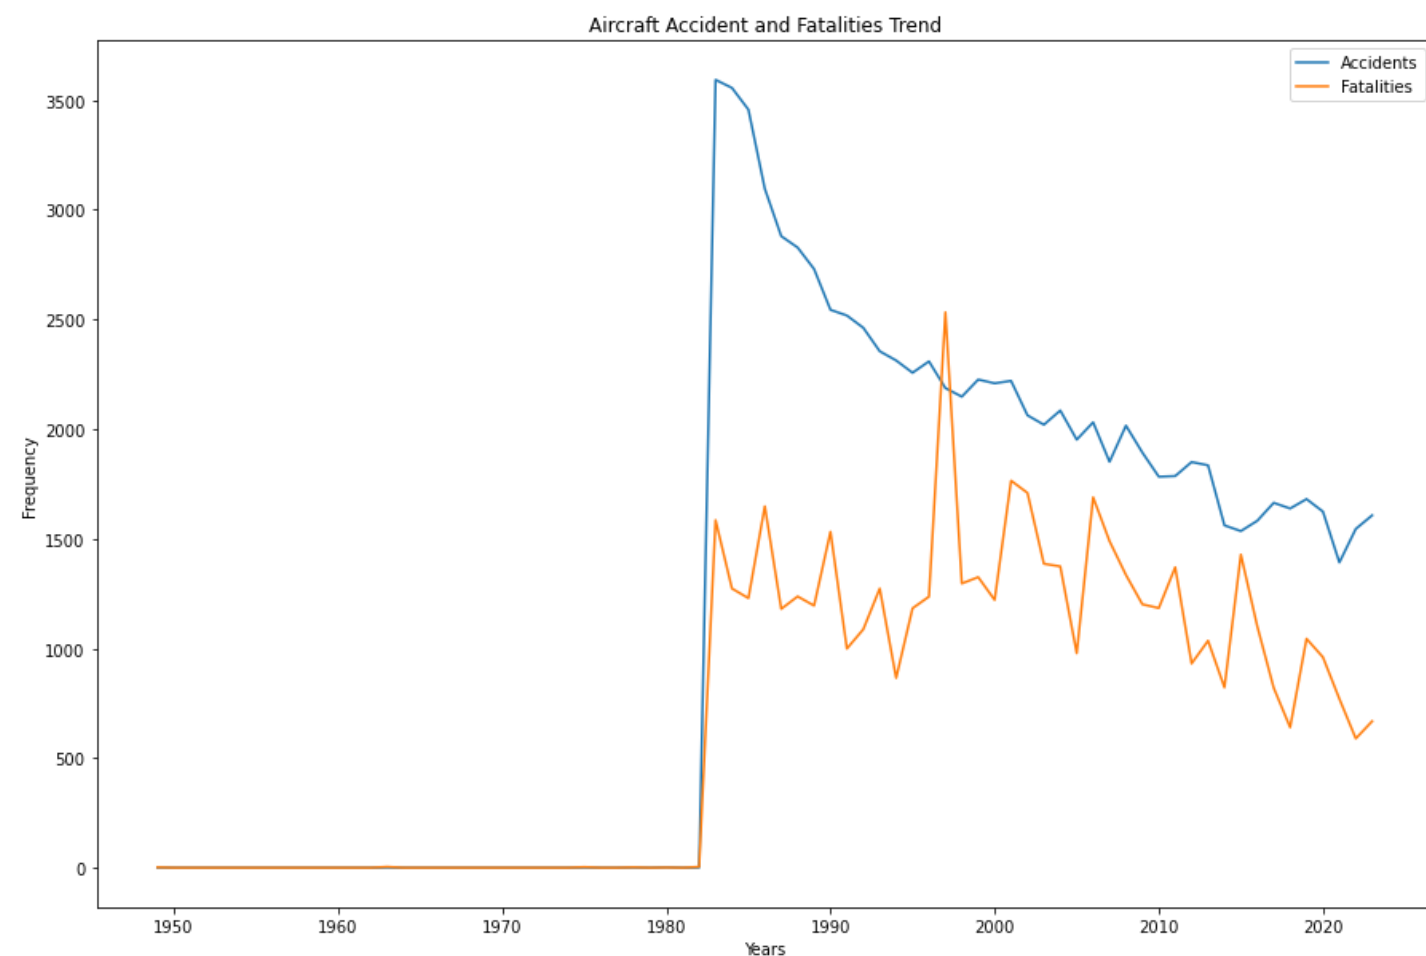
In [24]:

```
# Plotting a Line Graph to view the number of accidents and fatalities over the years
plt.figure(figsize=(15,10))

# Grouping the data by year
grouped_by_year = aviation_data.groupby(pd.Grouper(key='event_date', freq='Y')).agg({'investigation_type': 'count', 'total_fatal_injuries': 'sum'}).reset_index()

sns.lineplot(x='event_date', y='investigation_type', data=grouped_by_year, label='Accidents')
sns.lineplot(x='event_date', y='total_fatal_injuries', data=grouped_by_year, label='Fatalities')

plt.title('Aircraft Accident and Fatalities Trend')
plt.xlabel('Years')
plt.ylabel('Frequency')
plt.legend();
```



As indicated in the line graph above, the number of aircraft accidents continues to decrease over the years. Similarly, the number of fatal injuries has also reduced over the years.

Aircraft Accidents at Different Phases of Flight

The image below highlights different phases of flight for aircrafts. In this section, the total number of accidents that occurred in each phase of flight over the years were investigated. The number of fatal injuries and the level

of damage done to the aircraft at each phase of flight was also investigated.

The bar graph below showcases the total number of accidents that have occurred at each phase of flight.


In [25]:

```
# Plotting a bar graph to show number of accidents that have occurred in each phase
plt.figure(figsize=(12,6))

accidents_by_phase = aviation_data.groupby('broad_phase_of_flight').count()['investigation_type'].sort_values()

sns.barplot(x=accidents_by_phase.index, y=accidents_by_phase.values, color='lightseagreen')

plt.title('Number of Accidents per Phase of Flight')
plt.ylabel('Number of Accidents')
plt.xlabel('Phase of Flight');
```



From the bar chart above it seems that for most accidents in the dataset, the phase of flight in which the accident occurred was not known. However, this does not provide much information for the organization. For this reason, the focus shall be only on accidents whose phase of flight is known.

In [26]:

```
# Plotting a bar graph to show number of accidents that have occurred in each phase
plt.figure(figsize=(12,8))


# Filterin the data to only contain accidents whose phase of flight is known
accidents_at_unknown_phase = aviation_data.loc[aviation_data['broad_phase_of_flight']=='Unknown']
dropped_unknown_phase = aviation_data.drop(accidents_at_unknown_phase.index)
accidents_by_phase = dropped_unknown_phase.groupby('broad_phase_of_flight').count()['investigation_type'].sort_values()

sns.barplot(x=accidents_by_phase.index, y=accidents_by_phase.values, color='lightseagreen')

plt.title('Number of Accidents per Phase of Flight')
plt.ylabel('Number of Accidents')
plt.xlabel('Phase of Flight');
```

Number of Accidents per Phase of Flight

16000



It appears most accidents occur at the Landing and Takeoff phases. These are the points at which the aircraft is transitioning from air to ground and from ground to air respectively. During these transitions there is little time for error correction.

Now the bar chart below displays the number of fatal injuries recorded at each phase.

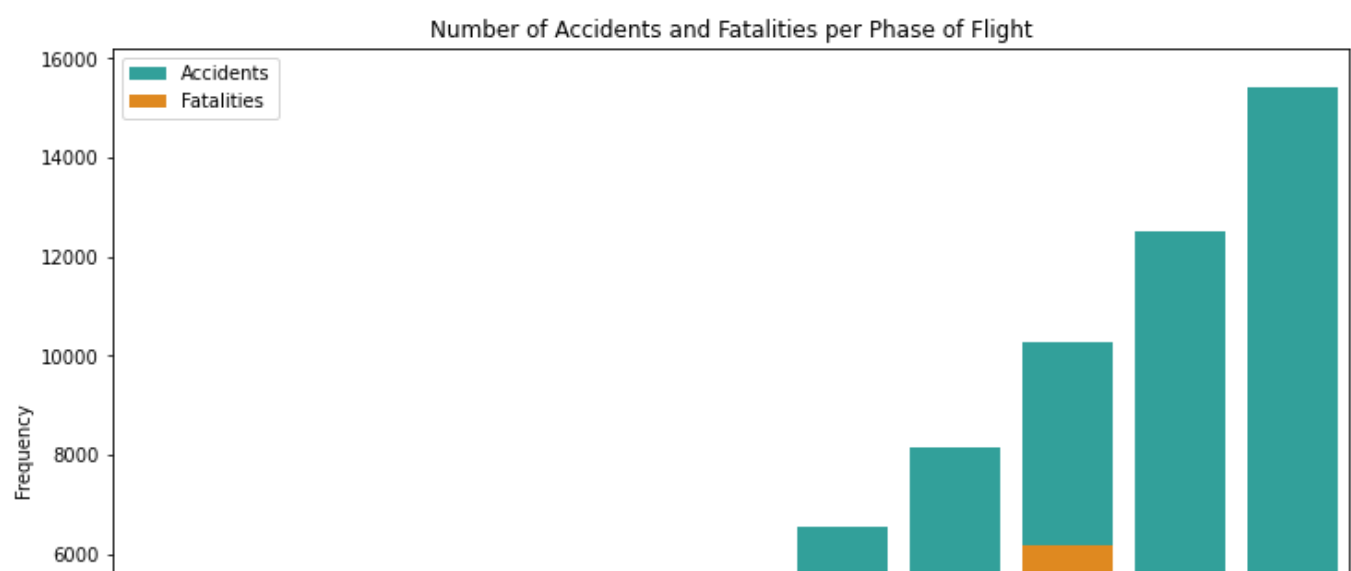
In [27]:


```
plt.figure(figsize=(12,8))

grouped_by_phase = dropped_unknown_phase.groupby('broad_phase_of_flight').agg({'investigation_type': 'count', 'total_fatal_injuries': 'sum'}).sort_values(by=['investigation_type'])

sns.barplot(x=grouped_by_phase.index, y='investigation_type', data=grouped_by_phase, color='lightseagreen', label='Accidents')
sns.barplot(x=grouped_by_phase.index, y='total_fatal_injuries', data=grouped_by_phase, color='darkorange', label='Fatalities')

plt.title('Number of Accidents and Fatalities per Phase of Flight')
plt.ylabel('Frequency')
plt.xlabel('Phase of Flight')
plt.legend();
```





Although the Landing and Takeoff phases of flight have the most occurrences of accidents, very few of them result in fatalities. The highest number of fatalities is instead experienced in accidents that occur in the cruise phase. Although cruise is generally a safer phase of flight compared to Landing and Takeoff, the risk of fatalities is much higher when accidents occur.

Finally, the chart below shows the level of damage done to the plane at each phase.


In [28]:

```
plt.figure(figsize=(15,8))

grouped_by_phase_and_damage = dropped_unknown_phase.groupby(['broad_phase_of_flight', 'aircraft_damage']).count().reset_index()

sns.barplot(x='broad_phase_of_flight', y='investigation_type', data=grouped_by_phase_and_damage, hue='aircraft_damage')

plt.title('Level of Aircraft Damage per Phase of Flight')
plt.ylabel('Number of Accidents')
plt.xlabel('Phase of Flight');
```



As noted earlier, most accidents occur at the Landing and Takeoff phase of flight. In addition to this, it seems that in the occurrence of these accidents, a substantial amount of damage is done to the aircraft. For the phases that have a higher risk of having fatalities (Cruise, Maneuvering), the aircrafts tend to be destroyed in the event of an accident.

Relation between number of engines in an aircraft and fatalities in the occurrence of an accident

Finally, an analysis was done to investigate if there is any relation between the number of engines in an aircraft and the fatalities in the event of an accident. The bar graph below shows the number of engines in an aircraft

and the number of fatal injuries for aircrafts with that number of engines.

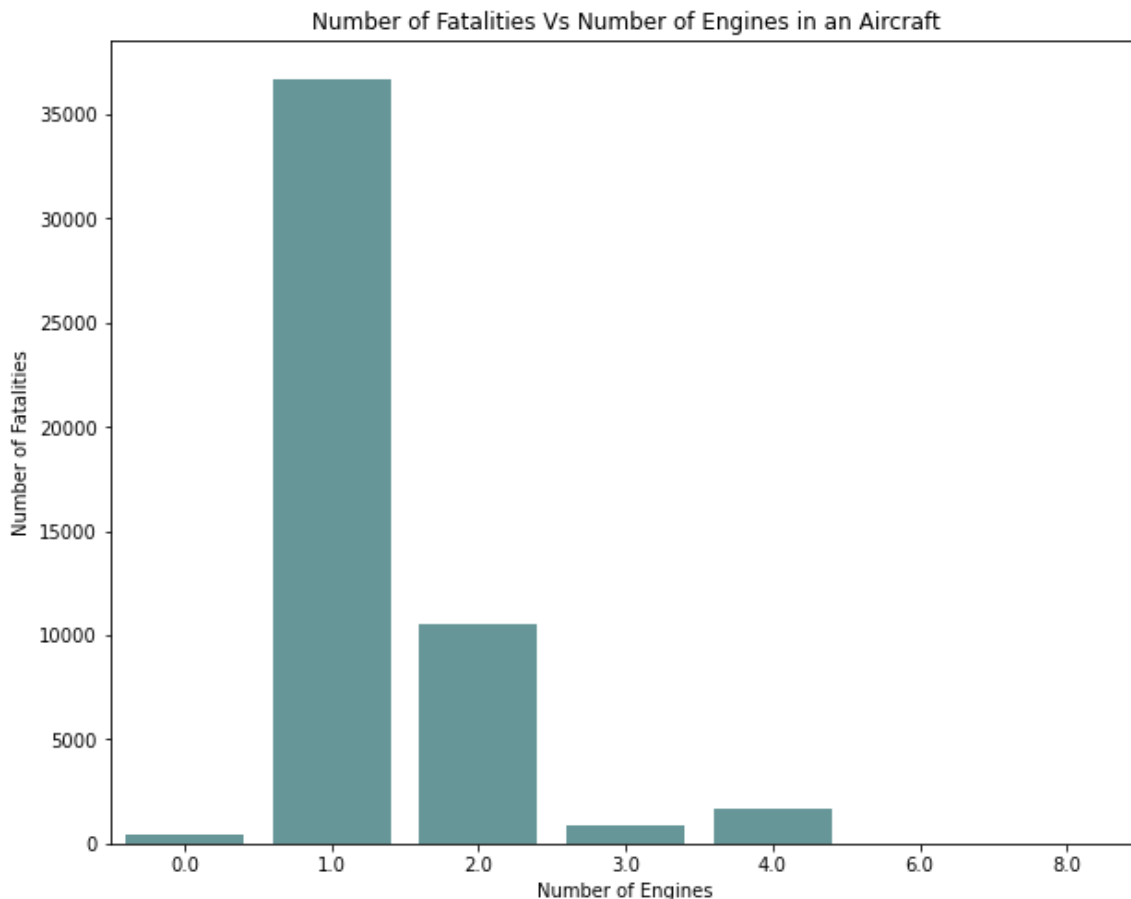
In [29]:

```
plt.figure(figsize=(10,8))

fatalities_vs_num_engines = aviation_data.groupby('number_of_engines').sum().reset_index()

sns.barplot(x='number_of_engines', y='total_fatal_injuries', data=fatalities_vs_num_engines, color='cadetblue')

plt.title('Number of Fatalities Vs Number of Engines in an Aircraft')
plt.xlabel('Number of Engines')
plt.ylabel('Number of Fatalities');
```



It can be observed that the highest number of fatalities occurs when the aircraft has only one engine.

Conclusion

It is said that airplanes are the safest mode of transportation. Indeed, it can be observed from this analysis that the number of aircraft accidents and fatalities have significantly reduced over the years.

However, it is important to note situations that lead to occurrences of these accidents. From this analysis, it's apparent that majority of aircraft accidents occur at the Landing phase of flight and the Takeoff phase of flight. This means that aircrafts are at the highest risk of crashing at the beginning and the end of a flight, when pilots have little time to react if something goes wrong. Additionally, aircrafts are also more likely to experience turbulence and bird strikes at these low altitudes. It is also in these phases of flight where the aircrafts are substantially damaged in the event of an accident. However, it is rather encouraging to note that the number of fatalities recorded in these two phases are quite low.

The analysis also shows that there is a higher risk of fatal injuries in accidents that involve aircrafts with only one engine while those with more than one engine have less fatal injuries.

In conclusion, the aircraft industry is a worthwhile business venture for the company given that a lot of effort has been put into advancing the technology used in airplanes to make them safer. Investing in aircrafts that have sophisticated auto-pilot and auto-landing systems will be beneficial as this will enhance flight precision and

reduce number of accidents at Landing and Takeoff. Finally, it is also advisable for the company to invest in aircrafts that have at least two engines, to further reduce the risk of fatalities.